

# Scalable, Secure and Broad-Spectrum Enforcement of Contracts—Without Blockchains

\*

1<sup>st</sup> Naftaly Minsky  
Computer-Science Department  
Rutgers University  
New Brunswick NJ, USA  
minsky@cs.rutgers.edu

2<sup>nd</sup> Chen Cong  
Computer-Science Department  
Rutgers University<sup>x</sup> (Now employed by Google LLC)  
New Brunswick, NJ, USA  
chen.cong@cs.rutgers.edu

**Abstract**—This paper introduces a scalable and secure contract-enforcement mechanism, called *Cop*, which can be applied to a broad range of multi-agent systems including small and large systems, time-critical systems, and systems-of-systems. *Cop* enforces contracts (or protocols) via the existing Law-Governed Interaction (LGI) mechanism, coupled with a new protective layer that significantly enhances the dependability and security of such enforcement.

*Cop* is arguably superior to the currently popular blockchain-based smart-contract mechanisms, due to its scalability, interoperability, and the breadth of the spectrum of its domain of applications.

**Index Terms**—distributed systems, enforcement of protocols, LGI, smart-contracts, dependability, security

## I. INTRODUCTION

There are many situations where a group of autonomous actors—which may be software processes, physical devices, and/or people operating via a platform like a smart-phone—need to interact with each other over the Internet. The members of such a group, who may not trust each other, may be required to interact subject to a given protocol. Such a protocol may represent a contract that binds these actors, or it may be necessary for the actors to collaborate effectively on some common goal, or to compete safely over the use of some resources. We call a group of actors that comply with a given protocol *P*, a *P-community*, or simply a community.

It is sometimes possible to establish a given protocol *P* over a community by relying on voluntary compliance with it, by all the members of this community. Voluntary compliance can be effective when a community is homogeneous, or when it is relatively small and its members trust each other, or when its members are well managed. But otherwise, for voluntary compliance with a given protocol *P* to be effective it needs to satisfy the following two conditions, as argued in [6]: (a) it must be the vested interest of every member of the community to comply with *P*; and (b) a failure to comply with *P*, by anybody, should not cause any serious harm to anybody else in the given community. And if any one of these conditions is not satisfied, then the given protocol may need to be *enforced*. The enforcement of protocols is the subject of this paper. And

we call a mechanism used for enforcing protocols a *protocol enforcement mechanism* or PEM.

We start in Section II by introducing the set of qualities of a PEM that we consider essential. We continue in Section III by reviewing two existing realizations of PEM: (a) the currently popular blockchain-based smart-contract mechanism, and (b) the older law-governed interaction (LGI) mechanism introduced by the 1st author—evaluating them in terms of these qualities. We find that each of these mechanisms fails to satisfy some of these qualities.

We obtain a PEM that satisfies all these qualities, by extending LGI into what we call *Cop*. Before introducing *Cop* we outline, in Section IV, the LGI mechanism itself, and describe its shortcoming—which we attempt to resolve. The LGI-based *Cop* mechanism is introduced in Section V, and discussed further in subsequent sections. We conclude in Section VIII, and with an appendix in Section A.

*a) A Terminological Comment:* In the computer-science literature the concept of protocol is sometimes referred to as a contract, and sometimes as a law. All three terms are being used here in the following way: We use the term “protocol” as a general term, when not discussing any particular enforcement mechanism; we use the term “contract” when discussing blockchain-based smart-contracts; and we use the term “law” when discussing LGI and *Cop*.

## II. THE ESSENTIAL QUALITIES OF PROTOCOL-ENFORCEMENT MECHANISMS

There is a wide range of types of applications that may benefit from the enforcement of protocols. Such applications differ along several dimensions. Some are small, others may be very large. Some are simple enough to be handled as a single community operating subject to a single protocol, others consist of multiple communities, operating subject to different protocols, which may need to inter-operate in various ways. Some applications are lax about the speed of protocol enforcement, others are time-critical. The actors involved in such applications may include software processes, people, and physical devices (i.e., devices that belong to IoT).

One would like to have a single protocol-enforcement mechanism (PEM) which is sufficiently broad-spectrum to support a wide range of potential applications. We identify below four qualities that a PEM should satisfy to be sufficiently broad-spectrum.

- 1) *Sufficiently short latency*: By latency we mean, the time it takes for a PEM to resolve a given transaction<sup>1</sup>. The maximal latency that a given application may require varies widely. It may, for example, be of the order of a few minutes for many commercial and financial systems. But it may be of the order of milliseconds or less for the so-called *time-critical* applications—such as a collection of interacting physical machines operating in an industrial plant, or the components of an airborne control system interacting with each other.
- 2) *Scalability*: By scalability we mean here that the latency is virtually independent of the volume and frequency of transactions in a given system. Scalability is a challenge to large systems, such as enterprise systems, and to infrastructures such as air-traffic control systems. It is also a challenge to financial systems such as the one described in [3], where “a reasonable estimate of [its] peak figure may be in the region of several thousand transactions per second.”
- 3) *Interoperability*: By interoperability we mean: the ability of communities that operate subject to different protocols to interact with each other. Of course, such an ability needs to be subject to regulation. That is, there needs to be a way to control which communities can interact with each other, and how.  
Interoperability is required, in many situations. In particular, when different small businesses, each operating under its own protocol, need to interact with each other. Moreover, complex systems, such as enterprise systems, cannot be governed by a single protocol. Rather, different communities of actors that belong to the same system, but engaged in different types of activities, would be required to operate subject to different protocols. And such communities often need to interact with each other. Therefore, a PEM needs to support multi-community systems, where different communities, operating under different protocols, need to be able to interact, subject to some constraints. Moreover, as we have shown [9], effective modularity of the set of protocols, that thus govern a complex system, can be achieved by organizing them into the so-called *conformance hierarchies*.
- 4) *Dependability and Security*: By this we mean the degree to which a PEM can defend itself against failures and programming errors (dependability) and against attacks (security). Dependability and security are, of course, critical for many applications.

It is worth pointing out that the satisfaction of any of these qualities is not a zero/one predicate, so our objective is a

*substantial* satisfaction of these qualities.

### III. ON THE STATE-OF-THE-ART OF PROTOCOL ENFORCEMENT MECHANISMS

In the following two sub-sections we consider two existing protocol-enforcement mechanisms mentioned in the introduction, evaluating them in terms of the qualities described above, and finding both of them wanting.

#### A. The Blockchain-based Smart-Contract Mechanism

This type of mechanisms—which was inspired by a 1997 paper [12] by Nick Szabo—became very popular recently, mostly for financial and commercial applications (see [11]). The main characteristic of smart-contracts is that the enforcement of contracts is carried out over a *blockchain*, which provides this mechanism with a high level of security.

But smart-contracts do not satisfy well the other three qualities listed above. First, the latency of a smart-contract cannot be shorter than the time it takes to reach consensus—a fundamental element of blockchains. And this latency is quite substantial—it is currently of the order of a few minutes under the various implementations of smart-contract, and it probably cannot be made much shorter than a few seconds. This means that smart-contracts *cannot be used for many time-critical applications*.

Second, the smart-contract mechanism is not scalable, as is frequently admitted by many researchers and developers of such systems. And although many, like [13], are working on reducing the level of unscalability, the lack of scalability is inherent in the blockchain-based smart-contract mechanisms. This is because, despite the distributed nature of the consensus, regarding which block to admit to the various copies of the blockchain, the enforcement itself is essentially centralized, and linear, for the following reason: A new block cannot be admitted to the blockchains, without sacrificing security, before the previous block is resolved. Now suppose that it takes  $T$  seconds to select a block (via a distributed consensus mechanism), to be admitted into the distributed blockchain, and to be resolved according to the contract at hand; and suppose that the average number of transaction that can be included in a block is  $B$ . Now, if more than  $B$  new transactions arrive, on average, to the blockchain in  $T$  seconds, then the length of the queue of transactions waiting to be processed will increase linearly in time. And the latency will increase, proportionally with the length of the queue. So, such a mechanism is inherently unscalable.

Third, blockchain-based smart-contracts cannot handle really complex systems, such as federated enterprises, supply chain, and health-care systems. Such systems are composed of many different communities operating under different interdependent contracts which often need to inter-operate. But despite some recent attempts to make blockchains interoperate, such as by the Cosmos project, none of them provides a practical and general solution to this problem. The Cosmos project, in particular, features a hub blockchain involved in all interoperations. And it seems to us that the use of such a

<sup>1</sup>We use the term “transaction” to mean any interactive operation by one of the actors in a given community.

hub would decrease the scalability and the security of smart-contracts.

### B. The Law-Governed Interaction (LGI) Mechanism:

This mechanism, which is discussed in some detail in Section IV, satisfies most of the qualities we required in Section II. This includes very short latency, high level of scalability, and a very general concept of controllable interoperability. But although LGI is reasonably secure—arguably more secure than the centralized access control mechanisms—LGI has a serious security weakness described in Section IV-F, which this paper aims to resolve. We sometime refer to this weakness of the security of LGI as its *Achilles heel*.

## IV. AN OUTLINE OF LGI, AND ITS *Achilles Heel*

This section is an outline of the LGI mechanism, which should suffice for the understanding of the rest of this paper. But for a deeper and more detailed description of LGI, we propose two, somewhat dated, sources: (a) a Journal paper [10]; and (b) the manual of LGI [8]. We confine ourselves here to the treatment of a single community of actors that are supposed to interact with each other subject to a common LGI-law  $\mathcal{L}$ —such a community is called an  $\mathcal{L}$ -community, and we occasionally refer to it as  $C$ .

The enforcement of an LGI-law is strictly *decentralized*, as follows: Each member  $a$  of  $C$  interacts with other members of  $C$  via a private surrogate (called a *controller*) that enforces law  $\mathcal{L}$  over the interactive activities of  $a$ . Such an enforcement is done *locally*, with no knowledge of, or dependency on, anything that happens simultaneously at the surrogates of other members of  $C$ . (This locality is the consequence of the nature of laws, as we shall see below). So, the enforcement of a law  $\mathcal{L}$  over the interactive activities of members of the  $\mathcal{L}$ -community is done in a decentralized manner, and thus in parallel, by the various surrogates of the members of  $C$ . This enforcement is very efficient, and inherently scalable.

The rest of this section is organized as follows: Section IV-A describes the service that provides actors with the controllers that can serve as their surrogates; Section IV-B describes how a given actor can become a member of the  $\mathcal{L}$ -community, for a given law  $\mathcal{L}$ ; Section IV-C outlines the structure of LGI-laws; Section IV-D introduces two simple examples of laws; Section IV-E points out that LGI can handle far more complex systems than is implied by the single community discussed so far. Finally, Section IV-F discusses the security vulnerability of LGI, which is what this paper is intended to relieve.

### A. A Trustworthy Controller-Service (CoS)

LGI require the availability of a set of authentic generic controllers that are trusted to operate as surrogates of arbitrary actors, subject to any well formed LGI-law selected by them. Such a generic controller is denoted by  $T$ , which suggests that it needs to be trusted to enforce correctly any LGI-law loaded into it. A controller operating under a given law  $\mathcal{L}$  is called an  $\mathcal{L}$ -controller, and is denoted by  $T^{\mathcal{L}}$ .




Fig. 1. A pair of interacting agents, operating under possible different laws

There are several ways for supplying such generic controllers. One of this is to create a *controller service* (CoS) that maintains set of authentic controllers, and leases them to its customers, presumably for a fee. The CoS can use various techniques for guarding against corruption of controllers. In particular it can use the TPM (Trusted Platform Module) technology, or some more recent variant of it. This is particularly easy to do because all generic controllers have identical codes, which would have a single and stable hash.

To be trustworthy, a CoS must be managed by a highly reputable organization, which can vouch for the authenticity of its controllers. Moreover, the CoS should provide each of its controllers with a digital certificate signed by the CoS itself. This certificate is used by the controllers to authenticate themselves as genuine LGI-controllers. The CoS is the trusted computing base (TCB) of LGI.

### B. The Concept of an $\mathcal{L}$ -Agent, and the Formation of an $\mathcal{L}$ -Community

1) *An  $\mathcal{L}$ -agent and its formation:* Any actor  $a$  can attempt to operate under a given law  $\mathcal{L}$ , thus joining the  $\mathcal{L}$ -community. It can do so by performing the following two steps: (i) acquiring a generic controller  $T$  from the CoS; and (ii) adopting this controller to be its surrogate, subject to law  $\mathcal{L}$ .

If the adoption of a controller  $T$  subject to law  $\mathcal{L}$  is successful, then this controller would be denoted by  $T_a^{\mathcal{L}}$ , which means that  $T$  now operates as the surrogate of  $a$ , subject to law  $\mathcal{L}$ . Now, the pair  $\langle a, T_a^{\mathcal{L}} \rangle$ , is called an  $\mathcal{L}$ -agent, which is denoted by  $\bar{a}$ . The mediator  $T_a^{\mathcal{L}}$  is called the *private controller* (or simply the controller) of agent  $a$ , and the actor  $a$ , whose internal structure is irrelevant to this model, is said to *animate* agent  $\bar{a}$ .

But note that the attempt of  $a$  to form an  $\mathcal{L}$ -agent may fail, because law  $\mathcal{L}$  may refuse to allow  $a$  to do so. For example, law  $\mathcal{L}$  may require a password, or a certificate, to be submitted by  $a$  in its adoption command. This would prevent any actor that cannot authenticate itself in the manner required by law  $\mathcal{L}$  from operating as a member of the  $\mathcal{L}$ -community.

2) *The Dual Mediation of Communication Under LGI:* One of the significant aspects of LGI is that it involves dual mediation of every exchange of messages between LGI-agents: one on the side of the sender of a message, and one on the side of its receiver. Specifically, the passage of a message from an actor  $A_x$  of an  $\mathcal{L}$ -agent  $x$  to an actor  $A_y$  of an  $\mathcal{L}'$ -agent  $y$ , must be mediated first by the controller  $T_x^{\mathcal{L}}$  associated with  $A_x$ , and then by the controller  $T_y^{\mathcal{L}'}$ , associated with  $A_y$ ,

as is illustrated in Figure 1. This is a direct consequence of the locality of LGI-laws, which requires both the sender and receiver to individually comply with the law under which each of them operates.

The dual mediation under LGI has several important implications, not the least of which is that it facilitates interoperability by providing flexible control over cross-interaction between agents operating under different laws, as is further discussed in Section IV-E. Moreover, as has been shown in [10], the dual control turns out to be more efficient than centralized control, in many circumstances. A simple illustration of the nature of dual mediation, and some of its consequences, is provided by the example laws on Section IV-D.

3) *On the Formation of an  $\mathcal{L}$ -Community*: An  $\mathcal{L}$ -community is created simply by the creation of the first  $\mathcal{L}$ -agent by some actor  $a$ . No special initialization of an  $\mathcal{L}$ -community is required. And the membership of such a community evolves incrementally, by the creation or destruction of  $\mathcal{L}$ -agents. The structure and dynamic behavior of such a community, which depends on the nature of its law, is explored in [1].

### C. The Concept of Law Under LGI

An LGI-law  $\mathcal{L}$  is formulated in terms of three elements, defined with respect to a given  $\mathcal{L}$ -controller  $x$ . The following is an incomplete description of these elements:

(1) A set  $E$  of *regulated events* (or, simply, *events*) that may occur at  $x$ .  $E$  includes, among others: (a) the *adoption* of this controller, under a given law  $\mathcal{L}$ —which is the first event that occurs at  $x$ . (b) The arrival of a message at  $x$ , and the sending of a message by it. Note that a message may arrive at  $x$  mostly from three sources: (i) from another controller; (ii) from the actor that adopted  $x$ ; and (iii) as a result from some *exception-condition*, which is reported to  $x$  via a message. And (3) an event called *obligationDue* that has to do with the pro-active capability of LGI. This type of events discussed in Section VII-B.

(2) The *state*  $S_x$  of controller  $x$ , which is distinct from the internal state of the actors that uses  $x$  as its surrogate—of which the law is oblivious. The state  $S_x$  is an unbounded (but often small) set of terms, whose structure is left unspecified here.

(3) A set  $O$  of *operations* that can be mandated by a law, to be carried out by the controller  $x$  upon the occurrence of a regulated event at it. The set  $O$  includes, among others: (a) replacing the state  $S_x$  of  $x$  with another state (or, if you will, changing the state  $S_x$ ); and (b) sending some messages, to anybody on the Internet. (Note that the sending of a message by  $x$ , as mandated by the law, would constitute a new event at  $x$ .)

Now, the role of a law is to decide what should be done in response to the occurrence of any regulated event at a controller operating under it. This decision, with respect to controller  $x$ , is formally defined by the following mapping:

$$E \times S_x \rightarrow (O)^* \times S_x \quad (1)$$

Using a less formal notation: the law, when applies to a given controller  $x$ , is a function

$$law : (e, s) \rightarrow (m, ns) \quad (2)$$

that maps any a given pair  $(e, s)$ , into a *ruling*  $(m, ns)$ . Here  $e$  is an event that occurred at controller  $x$ , and  $s$  is the state of  $x$  at the time of occurrence of this event. And the ruling that  $x$  is to carry out, consists of: (a) a (possibly empty) list  $m$  of *operations* that  $x$  must execute (which are, most often, messages to be sent); and (b) a new state  $ns$  that is to replace the current state of  $x$ .

This definition makes it clear that the law is strictly *local*. Indeed, the event that occurs at the controller and its state are local. And the ruling is to be carried out locally. Of course, if the ruling calls for some message to be sent to another controller, then this message would eventually have a non-local effect. But the decision to send such a message is local.

It should be pointing out that while Formula 1 is a definition of the semantics of laws<sup>2</sup>, it does not specify a language for writing them. In fact, the current implementation of LGI supports several *law-languages*, one of which is JavaScript. The choice of a law-language has no effect on the semantics of LGI, as long as the chosen language is sufficiently powerful to specify all possible mappings defined by Formula 1.

1) *Additional Observations about LGI-Laws and their Enforcement*:

a) *Non-Deterministic Laws*: A law may be *non-deterministic* in that its rulings may include random numbers. Such laws are useful, in particular, for the protocols involving randomness and tie-breaking. The implementation of non-determinism under Cop is discussed in [5].

b) *On the Interplay Between the Fixed Law and the Changing State of a Controller  $x$* : On one hand, the ruling of the law may depend on the current state of  $x$ , on the other hand the evolution of the state is regulated by the law—although it is driven by the various event that occur at  $x$ , most of them coming from other controllers in a given community.

c) *About the Enforcement of Laws*: A controller  $x$  deals with events occurring at it sequentially, and if several events occur at  $x$  at the same time, they will be handles in an arbitrary order. Also, the ruling of an event  $e$  is carried out atomically, before handling any subsequent event.

### D. Two Examples of Laws

We introduce here two very simple examples of LGI-laws, called *money transfer law* ( $\mathcal{MT}$ ) and *monitoring law* ( $\mathcal{MO}$ ). The former law is stated formally in one of the law-languages currently supported by LGI, and the latter law is described informally. We will return to these laws in Section IV-F, and also in Section V.

**A Money Transfer Law ( $\mathcal{MT}$ ):** This law provides an initial budget of \$1000 to every  $\mathcal{MT}$ -agent (this is done upon the adoption of a controller with law  $\mathcal{MT}$ ). And then,

<sup>2</sup>Modulo the fact that the sets  $E$  of events and  $O$  of operations have not been fully spelled out here.

this law enables every  $\mathcal{MT}$ -agent to transfer to others any amount of money smaller than or equal to its current budget. A formal statement of this law—written in the the law-language based on CoffeeScript (a semantically equivalent variant of JavaScript)—is spelled out in Box 1. This law has three rules, each of them contains comments (lines starting with #) that explain its effect.

**A Monitoring ( $\mathcal{MO}$ ) Law:** The following informally stated law, called  $\mathcal{MO}$ , establishes a systematic monitoring scheme for all communication within a given community.

- 1) *When any new  $\mathcal{MO}$ -agent is created—by an actor adopting a controller after inserting law  $\mathcal{MO}$  into it—the controller of the newborn agent would send a message to the designated monitor, essentially recording its own birth.*
- 2) *Whenever an  $\mathcal{MO}$ -message is sent, a copy of it, along with the addresses of the sender and its target, is sent to the monitor.*

Law 1. Money Transfer Law

```
Name: MT
LawScript: CoffeeScript

# (R1) When the controller is adopted,
# initialize the agent's budget to 1000.
UPON "adopted", ->
  DO "set", key: "budget", value: 1000
  return true

# (R2) An agent can send any positive amount of money
# to another agent provided that the amount
# is not greater than its budget,
# then the amount will be deducted from its budget.
UPON "sent", ->
  if @message > 0 and @message <= CS("budget")
    DO "set", key: "budget", value: CS("budget") - @message
    DO "forward"
    return true

# (R3) When an agent receives a positive amount of money,
# the amount will be deposited to its budget.
UPON "arrived", ->
  if @message > 0
    DO "set", key: "budget", value: CS("budget") + @message
    DO "deliver"
    return true
```

### E. Beyond Singleton Communities

So far we have discussed the case of an singleton  $\mathcal{L}$ -community, whose members interact with each other subject to a common law. But LGI is far more general than that, in the following ways, in particular: First, LGI can handle any number of communities, operating under different laws. Second, LGI can enable members of different such communities, say  $C1$  and  $C2$ —each operating under its own law—to interact with each other in a regulated manner. This can be done by having the laws of each of these community specify the condition under which its members can interact with each other. And third, LGI enables the organization of a set of laws that collectively governs a single system, into a coherent ensemble called a *conformance hierarchy*  $H$ .  $H$  is a tree of laws rooted by a law called  $\mathcal{L}_R$ . And every law in  $H$ , except of  $\mathcal{L}_R$  itself, conforms transitively to its superior law. Moreover the conformance relation between laws in  $H$  is inherent in the manner in which  $H$  is constructed, requiring no extra validation. For a formal definition of such an hierarchy

of laws see [2], and for a recent application of it to complex systems see [9].

### F. The Security Vulnerability of LGI

Even if the CoS of LGI does its utmost to maintain and protect authentic controllers there is, of course, no way to ensure that controllers cannot be corrupted and would violate the the law under which they are supposed to operate. Such a corruption may be the result of an attack on a controller, either by an insider of the CoS, or by an outsider who discovered some vulnerability in the code of controllers. We are concerned here mostly about the resulting *Byzantine behavior* of controllers, and not about their fail-stop type of failure, which can be handled effectively by LGI

The possible failure of individual controllers may be considered an acceptable risk in distributed computing, as it poses a smaller risk than that of the corruption of a central reference monitor commonly used in access control. Indeed, the corruption of a central reference monitor can endanger an entire system, while the corruption of a few controllers usually have a more local effect.

Yet, in some cases a Byzantine failure of even a single controller may cause a serious damage to the community at large. A case in point is the money-transfer law presented in Section IV-D. If a single controller has been corrupted, it may be able to distribute a large amount of fake money among other members of the  $\mathcal{MT}$ -community, without raising any suspicion, at least for a while. The *Achilles heel* of LGI is that it provide no general means for even detecting corrupt controllers.

Our approach for resolving this *Achilles heel*, thus protecting a system from the misbehaviors of its controllers, is the following: We provide a general mechanism that detects quickly and reliably any failed controller, right after it first failed to satisfy the law under which it operates; and then to recover from such a failure. This mechanism, called *detection & recovery* (or *D&R*), is the subject of the rest of this paper.

But first, we should make the following observation: There is, of course, a very general technique for handling Byzantine failures, see [4] for example. In principle, this techniques can be applied to every controllers of LGI. But this would be prohibitively too inefficient and expensive for most potential applications of LGI.

## V. THE LGI-BASED COP MECHANISM

Cop carries out two complementary functions in enforcing a given law  $\mathcal{L}$  over an  $\mathcal{L}$ -community  $C$ . One function is the enforcement, *per se*, of law  $\mathcal{L}$  over the interactive activities of the members of  $C$ . The other function is the speedy detection of any failure of an  $\mathcal{L}$ -controller to enforce law  $\mathcal{L}$ , followed immediately by the recovery from this failure, which includes the repair of the failed controller, and by the resumption of its operation.

These dual functions are carried out by two disjoint processes that operate in concert. One process is the enforcement of law  $\mathcal{L}$  over the  $\mathcal{L}$ -community, which is done by means of

the LGI mechanism, outlined in Section IV<sup>3</sup>, which provides LGI with an important proactive capability (cf. [5]). The other process, which operates off-line of the enforcement mechanism, is the detection of any misbehaving (or failed) controllers, and the recovery from such failures. This is done by a mechanism called D&R, for Detection and Recovery. The D&R part of Cop is the main subject of the rest of this paper.

It should be noted that—for simplicity—most of our description of Cop involves the treatment of a single isolated  $\mathcal{L}$ -community, whose members interact only with each other subject to law  $\mathcal{L}$ . But as we shall see, Cop can handle any number of such communities. Moreover, as explained in Section VII-A, Cop is not limited to dealing with isolated communities. Rather, like LGI itself, Cop can govern complex systems constituted of many interacting communities, inter-operating with each other subject to a conformance-hierarchy of laws.

We conclude this section with a description of the main components of Cop, and of the roles they play. The operations of the D&R mechanism, is discussed in Section VI.

#### A. The Components of the Cop Mechanism

The Cop mechanism is composed of three types of components that operate in concert.

- 1) The *controller provider* (CP), which is a variant of the CoS of LGI. Besides the maintenance of generic controllers, as does the CoS, the CP participates actively in the operations of the D&R mechanism.
- 2) The *ledger*  $D_L$  that maintains a record of the interactive behavior of all  $\mathcal{L}$ -controllers—i.e., the controllers that serve the  $\mathcal{L}$ -community.
- 3) The *inspector*  $I_L$ , which performs the inspection of the interactive behavior of all  $\mathcal{L}$ -controllers in order to detect any failure of any one of them, and to initiate the recovery of such failures.

Note that while the Cop mechanism has just one CP, which can handle any number of communities, each active  $\mathcal{L}$ -community is served by its own pair ( $D_L$ ,  $I_L$ ) of ledger and inspector, respectively. Note also, that the various ledgers and inspectors are to operate on different hosts, then the hosts used by CP. These three types of components are described in details in the following three subsections.

1) **The Controller Provider (CP):** The CP plays two kinds of roles. First, like the controller-service (CoS) of LGI, CP maintains generic controllers, and provides them to its clients. In that, CP is a more reliable version of the CoS, as argued below. Second, CP plays an important role in the operation of the D&R mechanism. These two roles are discussed in the following two paragraphs.

a) *The Maintenance of Generic Controllers:* The main structural difference between the CP and the CoS, in this respect, is that under the CP, controllers are encapsulated in Linux *containers* [7], hosted by a group of servers we call *CPnodes*, which are managed by the CP.

<sup>3</sup>There is just one difference between the LGI version used in Cop, and the older LGI—it is the implementation of the concept of *enforced obligation*

The CP can maintain any number of CPnodes, and each CPnode can host several hundreds of generic controllers, depending on the applications using them. And the controllers residing on any given CPnode may end up operating under different laws, and thus serving different communities. All the controllers resident in a given CPnode are managed by a *local manager* running on this CPnode—some of the functions of the local manager will be discussed in due course. The CP as a whole is managed by a *global manager* running on a distinguished CPnode of the CP. (And it would be useful, and quite elegant although not entirely necessary, for all these managers to interact with each other, subject to suitable *management-law* under LGI—this is not done in present prototype of Cop.)

The encapsulation of controllers in containers has several advantages. In particular, this architecture enables the imposition of limits on the use of various resources—such as CPU, memory and communication—by individual controllers. Such limits can be imposed by the local manager of the CPnode and dynamically adjusted by it. The ability to impose such limits would make controllers more robust because it can help to prevent an individual controller from hogging resources, thus preventing others from operating effectively, or at all.

b) *The Role that CP Plays in the Operations of the D&R Mechanism:* The CP carries out two functions that are essential to the D&R mechanism, as follows:

(b.1) *Supplying the ledger  $D_L$  with the information it needs:* As we shall see later, the detection of the failure of controllers by D&R requires that all the events that occur at every  $\mathcal{L}$ -controller, and all the operations carried out by it, be logged correctly on the ledger  $D_L$  (cf. Section V-A2). Unfortunately, we cannot trust the controllers themselves to log their own events and operations, because anyone of them may be corrupted. So, such logging is to be done by the local manager of every CPnode, by intercepting all the messages sent or received by every controllers that resides on the CPnode in question. And note that the sending or receiving of a message by a given controller corresponds to some event that occurred in it, and/or some operation carried out by it. The events and operations thus logged in ledgers are time-stamped, using the local time of the CPnode on which the controller resides; and they identify the controller in question.

Finally, it is important to point out that controllers operating subject to a given law  $\mathcal{L}$ , thus serving actors belonging to the  $\mathcal{L}$ -community, may reside on different CPnodes. The managers of these CPnodes would log events that occur of  $\mathcal{L}$ -controllers on the  $D_L$  ledger. Conversely, the manager of a given CPnode, which may host controllers operating under different laws, would have to log events on different ledgers.

(b.2) *Repairing failed controllers:* As discussed in Section VI-D1, the reconstruction of a failed  $\mathcal{L}$ -controller is carried out by the CP—following an instruction by the inspector  $I_L$ .

2) **The Ledger:** As has already been pointed out, there is one ledger  $D_L$  per an  $\mathcal{L}$ -community, which is designed to maintain entries representing two kind of items: (a) the *events*

# The Architecture of Cop

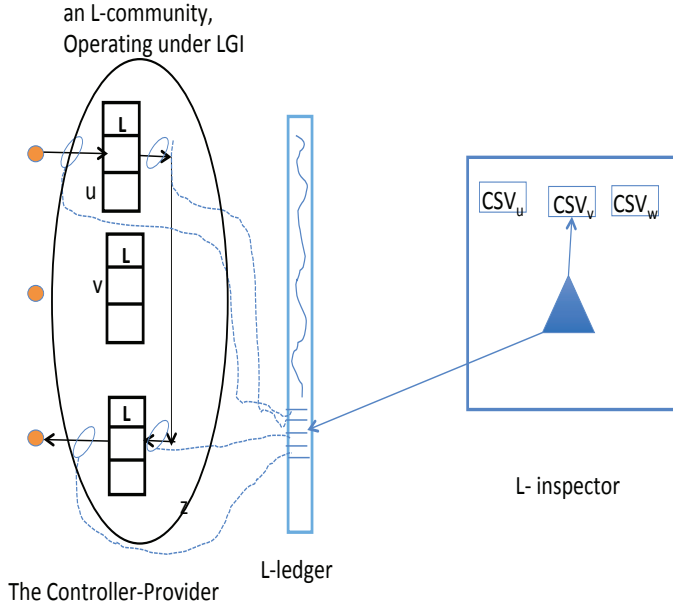


Fig. 2. A Schematic Depiction of the Operation of CoP

that occurred at any given  $\mathcal{L}$ -controller  $z$ , which represent messages obtained by  $z$  from various sources; and (b) the operations carried out by  $z$ , most of which represent messages sent by it to various targets. These entries are supplied to the ledger by the CP, as pointed out above.

Such a ledger can have various architectures, providing different level of security and efficiency to Cop. It can, in particular, be some form of blockchain, such as under Ethereum [11] or under HyperLedger [13]. Or it can be just a file, perhaps replicated, which a given organization maintains as part of its TCB. We will not discuss here the pros and cons of the various implementations of the ledger.

3) **The Inspector:** The objective of an inspector  $I_L$ , associated with a given  $\mathcal{L}$ -community, is twofold: (a) to inspect the interactive behavior of all the  $\mathcal{L}$ -controllers, in order to detect any failure of any of them; and (b) to initiate the recovery measures for failed controllers.

It is clear that the detection of a corrupt controller, and its recovery, needs to be done as quickly as possible—because the longer a corrupt controller is allowed to operate, the more damage it can do—damage that may be hard to reverse. The manner in which the inspector operates is discussed in Section VI.

4) **The Architecture of Cop:** Figure 2 provides a schematic depiction of the overall architecture and behavior of Cop. This figure depicts the treatment under Cop of the interaction between two actors operating under a law denoted by  $L$ , via a pair of controllers residing in a given controller-provider (CP).

The various events and operation generated by this interaction are intercepted by the the CP and sent to the L-ledger, which is inspected by the L-inspector.

## VI. INSPECTION: THE PROCESS OF DETECTION AND RECOVERY OF FAILED CONTROLLERS

This section starts with some introductory observations about the inspection process. Then, Section VI-B discusses the inspection of a single  $\mathcal{L}$ -controller, Section VI-C discusses the inspection of all the controllers serving a given  $\mathcal{L}$ -community, and Section VI-D focuses on the recovery from the failures of controllers.

### A. Introductory Observations

1) **Locality: the Key to Effective Detection of Failed Controllers:** The detection of the failure of controllers, by inspecting their interactive behavior, may seem to be a daunting and time consuming process, because the behavior of each controller depends on its interaction with others. This seems to suggest the need for global analysis of the process of interaction between all controllers for detecting any violation of the law by any of them.

Fortunately, no such global analysis is required, due to the inherently local nature of the laws of LGI (cf. Section IV). This aspect of laws under LGI is what enables it to enforce its laws locally, and thus efficiently and scalably. And since the enforcement is local, it follows that non-compliance with the law can also be dealt with locally—at each controller, independently of all other controllers. This simplifies the process of inspection enormously, and makes it very efficient and scalable. Thus, the locality of laws is the key for effective detection of the failure of controllers

2) **An Invariant of the Inspection Process:** The inspector  $I_L$  maintains a variable called  $CSV_x$ —which stands for *Controller State-Variable* of  $x$ —for every  $\mathcal{L}$ -controller  $x$  being inspected.  $CSV_x$  is computed, by the inspector, such that the following invariant is maintained:

The value of  $CSV_x$  is the *correct* state of the controller  $x$  associated with the current event that occurs at  $x$ .

By “correct state of the controller” we mean the state that a good controller  $x$  would have associated with the event that occurs at it. Of course, a failed, Byzantine, controller may have an arbitrary state, whose value cannot be predicted. We will see later how this invariant is maintained by the inspector.

3) **The Initialization of an  $\mathcal{L}$ -community Under Cop:** We described in Section IV the process of the creation and incremental development of an  $\mathcal{L}$ -community under LGI, which require no formal initialization. But under Cop, this process needs to be prefaced with the following initialization steps: (a) the formation of law  $\mathcal{L}$ ; (b) the creation of an empty ledger  $D_L$ ; and (c) the creation of inspector  $I_L$ , which is given the law under which it is to operate. Note that initially this community has no members, and no activity to be recorded on the ledger  $D_L$ . The inspector  $I_L$  starts examining the ledger, but doing

nothing else until it finds some entries on the ledger to be inspected.

### B. The Inspection of a single $\mathcal{L}$ -Controller

Broadly speaking, the process of inspection of a controller  $x$  by the inspector  $I_L$  starts when  $x$  becomes an  $\mathcal{L}$ -controller. This happens when  $x$  is adopted by some actor  $a$ , to operate subject to a law  $\mathcal{L}$ . This action by actor  $a$  triggers the so-called *adopted* event at  $x$ , which is the first event in its lifetime. When this event is intercepted by the CP and stored in the ledger  $D_L$ , it is observed by the inspector  $I_L$  and inspected, as described in Section VI-B1 below. After this initial inspection of  $x$ , its inspection continues recursively in response to consecutive events at  $x$  until  $x$  quits—as described in Section VI-B2.

For the detailed discussion of this process it would be helpful to recall the description of an LGI law provided by Formula 2, namely: a law is a function

$$law : (e, s) \rightarrow (m, ns) \quad (3)$$

that maps any a given pair  $(e, s)$ , into a *ruling*  $(m, ns)$ . Here  $e$  is an event that occurred in controller  $x$ ,  $s$  is the state of  $x$  at the time of occurrence of this event. And the ruling that  $x$  is to carry out consists of: (a) a (possibly empty) list  $m$  of *operations* that  $x$  is to perform (which are, most often, messages to be sent); and (b) a new state  $ns$  that is to replace the current state  $s$  of  $x$ .

1) *The initial Inspection of  $x$* : This inspection starts when the inspector  $I_L$ —which continuously scans the ledger  $D_L$  for new entries—notices the *adopted* event at  $x$  on the ledger, which prompts it to carry out the following sequence of steps.

(i) The inspector creates the variable  $CSV_x$  of  $x$ , whose value, in general, is to be the correct state of the controller  $x$  associated with any given event that occurs at  $x$ . The initial value of  $CSV_x$  is set to be an empty set, because, as stated in Section IV-B, this is the state of every newly formed controller, before its *adopted* event is evaluated.

(ii) The inspector computes the ruling of law  $\mathcal{L}$  for the pair  $(e, s)$ , where the event  $e$  is the adopted event and the state  $s$  is the current value of  $CSV_x$ , which is the empty state, as pointed up above. The ruling  $(m, ns)$  defines what a good  $x$  is expected to carry out, atomically.

(iii) Getting this ruling, the inspector changes the value of  $CSV_x$  to  $ns$ , which is clearly the correct state of  $x$  at this point, because it is mandated by the law.

(iv) The inspector will now verify if the controller  $x$  carried out the required list of operations in  $m$ —no more than in  $m$  and no less than in it. To do this, the inspector needs to compare  $m$  to the list  $m'$  of operations found on the ledger between the adopted event of  $x$  and its next event. (This must be the place for these operations, because the ruling of the law must be carried out atomically, and thus before the occurrence of next event at  $x$ .) Now, the inspector has two possibility to consider:

First, if  $m'$  equals to  $m$ , then the inspector concludes that  $x$  operates correctly.

Second, if  $m'$  differs from  $m$ , the inspector concludes that  $x$  just failed, and it will commence the appropriate recovery procedure, as described in Section VI-D.

2) *The Recursive Inspection of  $x$* : Suppose that inspector  $I_L$  inspected a sequence of events  $(e_1, e_2, \dots, e_n)$  that occurred at  $x$ , without detecting any failure—where  $e_1$  is the very first event that occurred at  $x$ , namely the *adopted* event discussed above. The inspection will continue, recursively, to the next event, if any, as follows:

(i)  $I_L$  will compute the ruling of the law for the pair  $(e_{n+1}, s_n)$ , where  $e_{n+1}$  is the new event at  $x$  found on the ledger, and  $s_n$  is the current value of  $CSV_x$  which is obtained during the previous step of this recursion. Suppose now that the ruling of the law for the pair  $(e_{n+1}, s_n)$  is the pair  $(m_{n+1}, ns_{n+1})$ , whose structure was described before.

(ii) Getting this ruling, the inspector plants  $ns_{n+1}$  as the new value of  $CSV_x$ .

(iii) The inspector will now check if  $x$  carried out the required list of operations  $m_{n+1}$ . To do this it needs compare this list to the list  $m'$  of operations found on the ledger immediately after the  $e_{n+1}$  event of  $x$ , and before the next event that occurred at  $x$ , if any. Now, the inspector has, again, two possibility to consider, in a direct analogy to the two possibilities it had in the very first inspection of  $x$ : First, if  $m'$  is equals to  $m_{n+1}$ , then the inspector concludes that  $x$  operates correctly. And second, if  $m'$  differs from  $m_{n+1}$ , then the inspector concludes that  $x$  just failed, and it will commence the recovery procedure from this failure, as described in Section VI-D.

It is worth pointing out, again, that the correctness of the value of  $CSV_x$  as the true state of the controller which should be used for the evaluation of the law, is an *invariant* of this recursive inspection.

### C. The Scalable Process of Inspecting of all $\mathcal{L}$ -Controllers

A single inspector can, in principle, inspect all the the  $\mathcal{L}$ -controllers that serve a given  $\mathcal{L}$ -community. Such an inspector would maintain the CSV of all active  $\mathcal{L}$ -controllers, in a single address space. And it would inspect all these controller virtually in parallel. But there is a potential problem with this modus operandi of inspection. Namely, when the size of the community increases, the latency between the failure of a controller and the detection of this failure grows, roughly linearly. In other words, such inspection is unscalable with respect to this latency. And the longer the latency is, the more opportunity a failed controller would have to send illegal messages that can cause serious damages to the system, and which may be very hard to recover from.

Fortunately, it is very easy to make the process of inspection scalable, as follows: If the set of  $\mathcal{L}$ -controllers in a given community is considered too large, in the sense that it produces overly large latency, one can divide this set to any number of subgroups that can be handled by different, but identical, copies of inspector that operate in parallel. Such division can be done dynamically, while the inspector operates.



#### D. Recovery from the Failures of Controllers

We consider here two complementary kinds of recoveries. The first is the resumption of the proper operation of a failed controller by its reconstruction. The second deals with the affect of the failed controller on other parts of the system.

1) *The Reconstruction of a Failed  $\mathcal{L}$ -Controller*: The reconstruction of a failed controller  $x$  is prompted by the inspector, and carried out by the local manager of the CPnode that host  $x$ . The reconstruction is carried out by the following sequence of steps:

(i) Controller  $x$  is replaced—without changing its address— with an authentic generic-controller provided by the CP. (ii) Law  $\mathcal{L}$  is planted into  $x$ , making it into an  $\mathcal{L}$ -controller. (iii) The latest value of  $CSV_x$  is planted into  $x$ . And, (iv) the reconstructed controller  $x$  is reactivated.

2) *Regarding the Affect of a Failed Controller on Other Parts of the System*: Note that an  $\mathcal{L}$ -controller is recognized by the inspector as failing, after it operated illegally, relative to law  $\mathcal{L}$ . Such an illegal operation needs to be corrected. Now, the ruling of a law is a list of zero or more operations that are to be carried out by the controller. The failure to carry out a given ruling consists of one or more instances of two types of illegality: (a) an *illegal inaction*, namely, the failure to carry out one of the operations in the ruling of law  $\mathcal{L}$ ; and (b) an *illegal action*, namely, carrying out an operation not in the ruling of law  $\mathcal{L}$ . These two types of illegalities require different handling, discussed below:

a) *The Handling of Illegal Inaction*: What one needs to do to recover from this kind of illegality is to carry out the operation required by the ruling. The Inspector does that right after discovering a failing controller, for every one of its illegal inactions, if any.

b) *The Handling of Illegal Action*: An illegal action means that the failing controller  $x$  sent some message to some controller  $y$ —an operation not mandated by the law. Such a message cannot be stopped, and its affect on  $y$ , and possibly on other members of the  $\mathcal{L}$ -community can be quite complex and possibly serious—as exemplified by the discussion of law  $\mathcal{MD}$  in Section IV-F. The recovery from such an illegal action may depend on the nature of the message sent, on the law, and on the application at hand, and is not a simple matter.

But Cop can help in recovering from such an illegal action, by having the inspector notify some designated manager—one associated with the  $\mathcal{L}$ -community—which can analyze the situation and decide how to rectify the problem. The operation of such a manager is beyond the scope of this paper. But it is worth pointing out that such a manager may need to examine parts of the  $D_L$  ledger, to determine the extent of the affect of the illegal message.

### VII. ADDITIONAL ASPECTS OF COP

#### A. The Controllable Interoperability Under Cop

As explained in Section IV-E, the LGI part of Cop—which is the part that enforces laws—provides for a sophisticated kind of interoperability. This capability of LGI extends to D&R, and thus to Cop.

Consequently, Cop can be used to govern the interaction between the members of disjoint communities. Moreover, can be used for governing very complex systems, such as federated enterprises, which comprises of a collection of communities operating under different laws that inter-operate with each other. Furthermore, it is possible to organize such inter-operating laws in a *conformance hierarchy*, which provides a way to control which communities can inter-operate with each other, and how. This provides Cop with considerable generality.

#### B. The LGI's Concept of Enforced Obligation, and its Treatment under Cop

The concept of *enforced obligation* (or “obligation” for short) provides LGI with an important *pro-active capabilities*, invaluable for security and for fault tolerance. This concept can be used, for example, to ensure that resources will not stay locked indefinitely, or to penalize book borrowers that do not return a book in the appointed time.

Informally speaking, an obligation *incurred* by a given controller, serves as a kind of *motive force*, which ensures that a certain action (called *sanction*) is carried out by this controller, at a specified time in the future (the deadline), when the obligation is said to *come due*—provided that certain conditions on the state of the controller are satisfied at that time. This mechanism is governed by the law in question.

Specifically, a controller  $x$  incurs an obligation by the execution, as part of the ruling of some event, of an operation **imposeObligation(oName,dt)**, where oName is the name of the obligation and dt is the time period after which the obligation is to come due. When this obligation comes due, after dt seconds, the event **obligationDue(oName)** would occur at controller  $x$ . The occurrence of this event would cause the controller to evaluate the ruling of the law for this event, and to carry out its ruling. The ruling of the law about an **obligationDue(oName)** event is, thus, the *sanction* for obligation oName.

But this concept of LGI, as is, cannot be supported under Cop. Because it relies on the controller itself to determine when the **obligationDue(oName)** event occurs. And our D&R mechanism does not, and cannot, rely on the controllers themselves. So, we implement the concept of obligation under Cop by making two changes to the obligation mechanism of LGI, without changing its semantics. First, the operation **imposeObligation(oName,dt)** causes a message with this text to be sent to an *obligation-server* implemented in the CPnode where the control  $x$  in question resides. When the deadline for this obligation arrives, the obligation-server is programmed to send the message **obligationDue(oName)** to  $x$ .

Second, the **obligationDue(oName)** message sent by the obligation-server, is viewed by the receiving controller as an event that would be handled just as the internal event **obligationDue(oName)** is handled by the original LGI. That is, the controller would evaluate the ruling of the law for this event, and then carry out its ruling—thus preserving the original semantics of the enforced obligation under LGI.

It should be pointed out that this is the only change in LGI we make in order to incorporate it in Cop.

### C. An Implementation of a Prototype of Cop, and its Testing

We have implemented a fully functional prototype of Cop, and tested it. The following is a summary of the results of this testing of the correctness and the performance of the D&R mechanism of Cop. (The correctness and performance of the LGI part of Cop has been tested many times in the past, and these measures did not change under Cop.)

To test the correctness of D&R—which is, in a sense, the test of the security of Cop—we’ve applied this prototype to multiple types of test cases, with different laws. We found *no false negatives*, i.e., all failed controllers were discovered. And we found *no false positive*, i.e., no correct controller were reported as a failed. Moreover, after the recovery of failed controllers, they all behaved correctly.

We have also measured the performance of the prototype of the D&R mechanism. We found (a) the latency of discovering a failed controller to be 6 seconds, on the average; and (b) the latency of the recovery of failed controllers to be 2 seconds, on the average.

But these performance measures have a limited value. First, because it is just a prototype which we tested, whose code is not optimized, in particular, it is written in Python 3—a scripting language with relatively low performance. And second, because our experiments were done on a relatively weak hardware, which did not allow us to experiment with large communities. We expect that our latency result would be reduced by an order of magnitude, with optimized code, run on a stronger hardware.

### D. On the Trustworthiness and Security of the D&R Mechanism

The D&R mechanism has been designed for protecting against the possible failure of controllers—which we call the Achilles heel of LGI. But this protection can be effective only if the D&R Mechanism itself—consisting of the CP, the ledgers, and the Inspectors—is secure. That is, these three components must constitute the trusted computing base (TCB) of D&R. We believe that our current design of these components of D&R of makes these components reasonably secure. And their security can be enhanced via various traditional means, including TPM and related technologies. Such security enhancements are beyond the scope of this work.

## VIII. CONCLUSION

This paper introduces a scalable and secure protocol-enforcement mechanism, called Cop, which fulfills important qualities for such mechanisms, including low latency, high scalability, general interoperability and security. It is thus applicable to a wide range of applications, including small and large systems, time-critical systems, and systems-of-systems.

Cop enforces protocols via the existing Law-Governed Interaction (LGI) mechanism, coupled with a new protective layer called D&R that discovers any failed LGI-controller and

repairs it—which is done soon after the failure occurs. The D&R layer of Cop operates off-line relative to the enforcement by LGI, and it significantly enhances the dependability and security of the enforcement.

We have implemented a fully functional prototype of Cop, and verified experimentally its correctness. But the evaluation of the performance and security of Cop, particularly when it is applied to large scale systems, would require an optimized implementation of Cop, and a sufficiently powerful hardware for it to run on.

## APPENDIX

### SUBSTANTIATING THE MAIN CLAIMS OF THIS PAPER

We have created a fully functional prototype of Cop, and tested it (cf. Section VII-C). This prototype will be made available on the authors’ websites.

## REFERENCES

- [1] X. Ao, N. Minsky, T. Nguyen, and V. Ungureanu. Law-governed communities over the internet. In *Proc. of Fourth International Conference on Coordination Models and Languages; Limassol, Cyprus; LNCS 1906*, pages 133–147, September 2000.
- [2] Xuhui Ao and Naftaly H Minsky. Flexible regulation of distributed coalitions. In *Computer Security—ESORICS 2003*, pages 39–60. Springer, unknown, 2003.
- [3] S Blakstad and R Allen. Fintech revolution, universal inclusion in the new financial ecosystem. cham, switzerland: Palgrave macmillan, part of springer nature. Technical report, Springer, September 2018.
- [4] Miguel Castro and Barbara Liskov. Proactive recovery in a byzantine-fault-tolerant system. In *osdi00*, October 2000.
- [5] Chen Cong. *Cop: broad-spectrum, dependable, and secure protocol enforcement mechanism for multi-agent systems*. PhD thesis, Rutgers, the State University of New Jersey, 2019.
- [6] M. Ionescu, Naftaly H. Minsky, and T. Nguyen. Enforcement of communal policies for peer-to-peer systems. In *Proc. of the Sixth International Conference on Coordination Models and Languages, Pisa Italy*, February 2004.
- [7] Dirk Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux Journal*, 2014(239):2, 2014.
- [8] Naftaly H. Minsky. *Law Governed Interaction (LGI): A Distributed Coordination and Control Mechanism (An Introduction, and a Reference Manual)*. Rutgers, February 2006. (available at <http://www.moses.rutgers.edu/>).
- [9] Naftaly H Minsky. On the dependability of highly heterogeneous and open distributed systems. *Journal of Software Engineering and Applications*, 11(01):28, 2018.
- [10] Naftaly H. Minsky and V. Ungureanu. Law-governed interaction: a coordination and control mechanism for heterogeneous distributed systems. *TOSEM, ACM Transactions on Software Engineering and Methodology*, 9(3):273–305, July 2000.
- [11] Gareth W Peters and Efstathios Panayi. Understanding modern banking ledgers through blockchain technologies: Future of transaction processing and smart contracts on the internet of money. In *Banking Beyond Banks and Money*, pages 239–278. Springer, 2016.
- [12] Nick Szabo. Formalizing and securing relationships on public networks. *First Monday*, 2(9), 1997.
- [13] Marko Vukolić. Rethinking permissioned blockchains. In *Proceedings of the ACM Workshop on Blockchain, Cryptocurrencies and Contracts*, pages 3–7. ACM, 2017.